Model Development Phase Template

Date	26 November 2024
Team ID	SWTID1726490119
Project Title	Toxic Comment Classification for Social Media.
Maximum Marks	10 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be displayed via a screenshot, showcasing the architecture and training process for the selected model.The model validation and evaluation report will summarize the training and validation performance metrics (accuracy, loss, etc.) for multiple models, with respective screenshots of evaluation results like classification reports and accuracy plots.

Initial Model Training Code (5 marks):

1. Simple Neural Network (ANN)

```
#define ANN model
def ann_model(hp):
    model = Sequential()
    # Tune the embedding dimensions
    embedding_dim = hp.Int('embedding_dim', min_value=64, max_value=256, step=64)
    model.add(Embedding(input_dim=20000, output_dim=embedding_dim))

    model.add(Flatten())

    # Tune the number of units in the Dense layer
    dense_units = hp.Int('dense_units', min_value=32, max_value=128, step=32)
    model.add(Dense(dense_units, activation='relu'))

    # Output layer for binary classification
    model.add(Dense(6, activation='sigmoid'))

    # Tune the learning rate for the Adam optimizer
    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),  # Adjust learning rate
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

    return model
```

```
[ ]  # Instantiate the tuner
     tuner = kt.RandomSearch(
         ann_model,
         objective='val_accuracy',
         max_trials=10,  # Number of different models to try
         executions_per_trial=1,  # Each model will be trained twice
         directory='ann_tuning',
         project_name='ann_tuning'
     )
```

```
[ ]  # Perform hyperparameter tuning
     tuner.search(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

```
[ ]  # Get the best model
     ann_best_model = tuner.get_best_models(num_models=1)[0]
     ann_best_hyperparameters = tuner.get_best_hyperparameters(num_trials=1)[0]
     print("Best Hyperparameters:")
     print(ann_best_hyperparameters.values)
```

```
     # Train the best model
     ann_best_model.fit(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

2. Convulational Neural Network (CNN)

```
     # CNN Model
     def cnn_model(hp):
         model = Sequential()
         model.add(Embedding(input_dim=20000, output_dim=128, input_length=maxlen))

         # Tune the number of filters and kernel size
         for i in range(hp.Int('num_conv_layers', 1, 3)):  # 1 to 3 convolutional layers
             model.add(Conv1D(filters=hp.Int('filters_' + str(i), 32, 256, step=32),  # 32 to 256 filters
                              kernel_size=hp.Int('kernel_size_' + str(i), 3, 7),  # Kernel size 3 to 7
                              activation='relu'))
             model.add(MaxPooling1D(pool_size=2))

         model.add(Flatten())

         # Tune the number of neurons in the dense layer
         model.add(Dense(hp.Int('dense_units', 32, 128, step=32), activation='relu'))

         # Add dropout for regularization
         model.add(Dropout(hp.Float('dropout_rate', 0.0, 0.5, step=0.1)))

         model.add(Dense(6, activation='sigmoid'))

         # Compile the model
         model.compile(optimizer=tf.keras.optimizers.Adam(
             hp.Float('learning_rate', 1e-4, 1e-2, sampling='LOG')),
                       loss='binary_crossentropy', metrics=['accuracy'])

         return model
```

```
[ ]  # Set up the tuner
     tuner = kt.RandomSearch(
         cnn_model,
         objective='val_accuracy',
         max_trials=10,  # Number of different hyperparameter combinations to try
         executions_per_trial=1,  # Number of times to train each model
         directory='cnn_tuning',  # Directory to store results
         project_name='cnn_tuning'
     )
```

```
[ ]  # Start the hyperparameter search
     tuner.search(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

```
[ ]  # Get the best model
     cnn_best_model = tuner.get_best_models(num_models=1)[0]
     cnn_best_hyperparameters = tuner.get_best_hyperparameters(num_trials=1)[0]
     print("Best Hyperparameters: ", cnn_best_hyperparameters.values)
```

```
     #Train the best model
     cnn_best_model.fit(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

3. Long Short-Term Memory (LSTM)

```
[ ]  # LSTM Model
     def lstm_model(hp):
         model = Sequential()
         model.add(Embedding(input_dim=20000, output_dim=128, input_length=maxlen))

         # Tune the number of LSTM units
         model.add(LSTM(hp.Int('lstm_units', 32, 256, step=32), return_sequences=False))  # LSTM units

         # Add dropout for regularization
         model.add(Dropout(hp.Float('dropout_rate', 0.0, 0.5, step=0.1)))

         model.add(Dense(hp.Int('dense_units', 32, 128, step=32), activation='relu'))  # Dense units
         model.add(Dense(6, activation='sigmoid'))  # Output layer

         # Compile the model
         model.compile(optimizer=tf.keras.optimizers.Adam(
             hp.Float('learning_rate', 1e-4, 1e-2, sampling='LOG')),
                       loss='binary_crossentropy', metrics=['accuracy'])

         return model
```

```
# Set up the tuner
tuner = kt.RandomSearch(
    lstm_model,
    objective='val_accuracy',
    max_trials=10,  # Number of different hyperparameter combinations to try
    executions_per_trial=1,  # Number of times to train each model
    directory='lstm_tuning',  # Directory to store results
    project_name='lstm_tuning'
)
```

```
[ ] # Start the hyperparameter search
    tuner.search(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

```
[ ] # Get the best model
    best_lstm_model = tuner.get_best_models(num_models=1)[0]
    best_lstm_hyperparameters = tuner.get_best_hyperparameters(num_trials=1)[0]

    print("Best Hyperparameters: ", best_lstm_hyperparameters.values)
```

4. Bi-Directional LSTM (BiLSTM)

```
# BiLSTM Model
def bilstm_model(hp):
    model = Sequential()
    model.add(Embedding(input_dim=20000, output_dim=128, input_length=maxlen))

    # Tune the number of LSTM units
    model.add(Bidirectional(LSTM(hp.Int('lstm_units', 32, 256, step=32), return_sequences=False)))  # BiLSTM units

    # Add dropout for regularization
    model.add(Dropout(hp.Float('dropout_rate', 0.0, 0.5, step=0.1)))

    model.add(Dense(hp.Int('dense_units', 32, 128, step=32), activation='relu'))  # Dense units
    model.add(Dense(6, activation='sigmoid'))  # Output layer

    # Compile the model
    model.compile(optimizer=tf.keras.optimizers.Adam(
        hp.Float('learning_rate', 1e-4, 1e-2, sampling='LOG')),
                  loss='binary_crossentropy', metrics=['accuracy'])

    return model
```

```
[ ]  # Set up the tuner
     tuner = kt.RandomSearch(
         bilstm_model,
         objective='val_accuracy',
         max_trials=10,  # Number of different hyperparameter combinations to try
         executions_per_trial=1,  # Number of times to train each model
         directory='bilstm_tuning',  # Directory to store results
         project_name='bilstm_tuning'
     )
```

```
     # Start the hyperparameter search
     tuner.search(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

```
[ ]  # Get the best model
     best_bilstm_model = tuner.get_best_models(num_models=1)[0]
     best_bilstm_hyperparameters = tuner.get_best_hyperparameters(num_trials=1)[0]

     print("Best Hyperparameters: ", best_bilstm_hyperparameters.values)
```

```
     best_bilstm_model.fit(X_train, y_train, epochs=5, batch_size=64, validation_data=(X_val, y_val))
```

Model Validation and Evaluation Report (5 marks):

Model	Summary	Training and Validation Performance Metrics
Simple Neural Network (ANN)	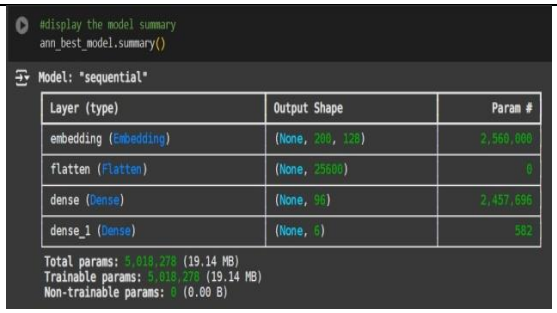	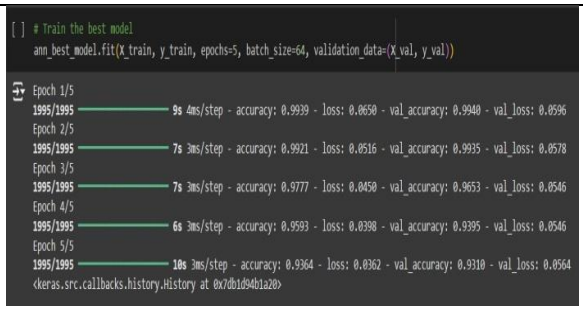

Convulational Neural Network (CNN)	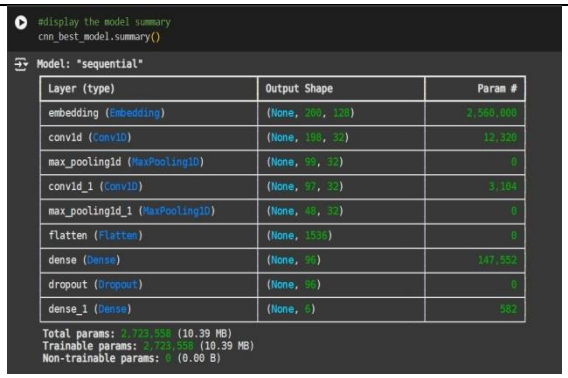	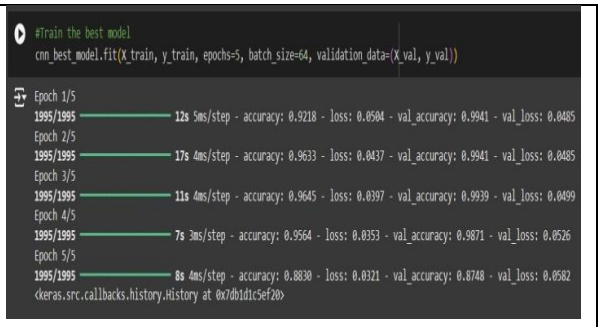
Long Short-Term Memory (LSTM)	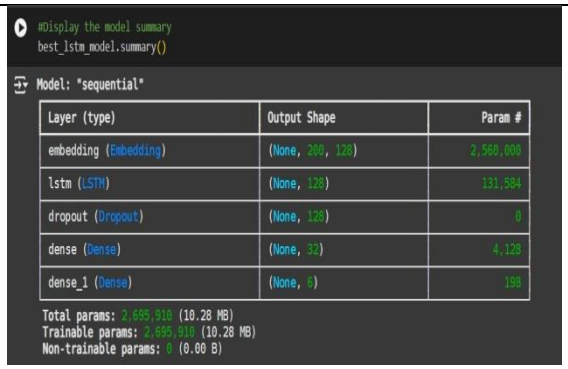	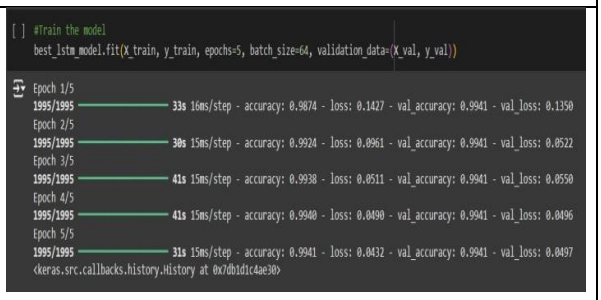
Bi-Directional LSTM (BiLSTM)